


FLIGHT DELAY PREDICTION MODELING

Flight delay prediction using Random Forest on synthetic and real data, exploring accuracy, key features, and class imbalance




MODEL 1: SYNTHETIC

99% acc

Top features:

- arrival time
- previous delays



MODEL 2: REAL

92% acc

Top features:

- arrival time
- departure time



After my family was stranded during a London transit meltdown, I wanted to engineer a predictive system that could reduce disruptions like this.

AI Model 1 – Synthetic Data : Creating the Artificial Data Set

```
[2]: np.random.seed(42)

num_samples = 50000 # Data size

df_synthetic = pd.DataFrame({
    "FL_DATE": pd.date_range(start="2024-01-01", periods=num_samples, freq="D"),
    "AIRLINE": np.random.choice(["Delta", "American", "United", "Southwest", "Alaska", "JetBlue", "Spirit", "Frontier"], num_samples),
    "ORIGIN": np.random.choice(["JFK", "LAX", "ORD", "ATL", "SEA", "DFW", "MIA", "SFO", "DEN", "PHX"], num_samples),
    "DEST": np.random.choice(["BOS", "LAS", "IAH", "MSP", "CLT", "DTW", "EWR", "FLL", "BWI", "SLC"], num_samples),
    "CRS_DEP_TIME": np.random.randint(500, 2359, num_samples),
    "NUM_PREVIOUS_FLIGHTS_LATE": np.random.randint(0, 10, num_samples), # More variability
    "AVG_GATE_WAIT_TIME": np.random.uniform(0, 60, num_samples),
})

# Simulate delays based on broader conditions
df_synthetic["DELAYED"] = (
    (df_synthetic["NUM_PREVIOUS_FLIGHTS_LATE"] >= 5) |
    (df_synthetic["AVG_GATE_WAIT_TIME"] > 50)
).astype(int)

# Save for future use
df_synthetic.to_csv("expanded_flight_delays.csv", index=False)

print("New Dataset Size:", df_synthetic.shape)
```

New Dataset Size: (50000, 8)

Output Results

AI Model 1 – Synthetic Data : Data Set

expanded_flight_delays.csv

C: > my_github_repos > Flight-Delay-Prediction-Modeling > expanded_flight_delays.csv > data

```
1  FL_DATE,AIRLINE,ORIGIN,DEST,CRS_DEP_TIME,NUM_PREVIOUS_FLIGHTS_LATE,AVG_GATE_WAIT_TIME,DELAYED
2  2024-01-01,Spirit,PHX,CLT,917,9,7.478868696583549,1
3  2024-01-02,Southwest,PHX,IAH,1345,8,49.40366927292514,1
4  2024-01-03,Alaska,JFK,MSP,1981,3,26.901369518243758,0
5  2024-01-04,Spirit,ORD,EWR,2330,1,46.92140320678622,0
6  2024-01-05,United,ATL,IAH,2285,3,39.61700657449153,0
7  2024-01-06,Frontier,MIA,BOS,1352,5,48.41399442100398,1
8  2024-01-07,Alaska,DEN,BOS,1201,2,34.008907082564754,0
9  2024-01-08,Alaska,SEA,MSP,1963,0,29.368711297202715,0
10 2024-01-09,Spirit,JFK,BOS,1386,8,59.708911934701675,1
11 2024-01-10,American,LAX,DTW,1804,6,0.6857796628601309,1
12 2024-01-11,United,PHX,MSP,2018,1,32.36797136871213,0
13 2024-01-12,Spirit,LAX,CLT,1944,8,54.198175226692946,1
14 2024-01-13,United,DEN,DTW,1195,3,2.527770840189254,0
15 2024-01-14,United,PHX,LAS,538,5,47.04697676305268,1
16 2024-01-15,Frontier,MIA,BOS,2034,7,19.59865872678699,1
17 2024-01-16,Alaska,LAX,SLC,1217,3,33.84299776359828,0
18 2024-01-17,Southwest,JFK,LAS,824,4,54.001670272365274,1
19 2024-01-18,Frontier,SEA,DTW,716,6,41.1699310043518,1
20 2024-01-19,Frontier,JFK,SLC,2080,4,7.671910249935614,0
21 2024-01-20,United,DEN,BWI,1845,3,22.73021841835738,0
22 2024-01-21,JetBlue,ORD,FLL,1123,8,34.35691656051302,1
23 2024-01-22,Alaska,LAX,SLC,647,6,26.897561787442495,1
24 2024-01-23,American,DFW,BOS,1405,1,16.303029373108217,0
25 2024-01-24,Frontier,ORD,CLT,1115,0,36.76633961378634,0
```

50,000 rows

AI Model 1 – Synthetic Data : Balancing Out the Synthetic Data Set

```
[5]: smote_tomek = SMOTETomek(random_state=42)
X_train_balanced, y_train_balanced = smote_tomek.fit_resample(X_train, y_train)

print("Balanced class distribution:", np.bincount(y_train_balanced))
```

Balanced class distribution: [17031 17031]

Output Results

AI Model 1 – Synthetic Data : Fitting the RandomForestClassifier Model to the Data

```
[6]: model = RandomForestClassifier(n_estimators=200, class_weight="balanced_subsample", random_state=42)
model.fit(X_train_balanced, y_train_balanced)
```

```
[6]: ▼ RandomForestClassifier
RandomForestClassifier(class_weight='balanced_subsample', n_estimators=200,
random_state=42)
```

Output Results

AI Model 1 – Synthetic Data : Confusion Matrix and Accuracy

```
[7]: y_pred = model.predict(X_test)

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Classification Report:\n", classification_report(y_test, y_pred))


# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

Accuracy: 0.9998

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	4185
1	1.00	1.00	1.00	5815
accuracy			1.00	10000
macro avg	1.00	1.00	1.00	10000
weighted avg	1.00	1.00	1.00	10000

Output Results



Only two flights were predicted wrong

AI Model 1 – Synthetic Data : F1 Score

```
[9]: F1 = f1_score(y_test, y_pred, average = 'macro')  
print(F1)
```

0.9997945273932791

Output Results

AI Model 1 – Synthetic Data : Feature Importance

```
[10]: # Create a series with feature names and their importance scores  
importances = pd.Series(model.feature_importances_, index=X_train.columns)  
  
# Sort and show the top 20 features  
top_features = importances.sort_values(ascending=False).head(20)  
print(top_features)
```

NUM_PREVIOUS_FLIGHTS_LATE	0.732326
AVG_GATE_WAIT_TIME	0.208496
FL_DATE	0.010131
HOUR	0.006916
AIRLINE_Frontier	0.003505
AIRLINE_Delta	0.003491
ORIGIN_ATL	0.002148
DEST_IAH	0.002068
ORIGIN_PHX	0.002025
AIRLINE_United	0.001991
AIRLINE_Southwest	0.001810
AIRLINE_Alaska	0.001808
DEST_SLC	0.001743
AIRLINE_JetBlue	0.001670
AIRLINE_American	0.001561
DEST_EWR	0.001515
ORIGIN_DEN	0.001383
AIRLINE_Spirit	0.001372
ORIGIN_JFK	0.001232
DEST_MSP	0.001182

dtype: float64

Top Features

AI Model 2 – Real-World Data : Preprocessing Data

Imported and cleaned 1.6M real flight records, filling missing values

```
12 # Load dataset
13 df = pd.read_csv('flights_sample_3m.csv')
14
15 # Fill missing categorical columns with mode
16 for col in df.select_dtypes(include='object').columns:
17     df[col] = df[col].fillna(df[col].mode()[0])
18
19 # Fill missing numeric columns with median
20 for col in df.select_dtypes(include=['float64', 'int64']).columns:
21     df[col] = df[col].fillna(df[col].median())
22
23 # --- Encoding ---
24
25 le = LabelEncoder()
26 for col in df.select_dtypes(include='object').columns:
27     if col != 'status':
28         df[col] = le.fit_transform(df[col])
29
30 # Create binary target column for delay: 1 if delayed, 0 otherwise
31 df['status_Delayed'] = (df['status'] == 'Delayed').astype(int)
32
33 # --- Prepare features and target ---
34
35 x = df.drop(columns=['status', 'status_Delayed'])
36 y = df['status_Delayed']
```


AI Model 2 – Real-World Data : Data Set

```
flights_sample_3m.csv > data
1 flight_id,flight_no,scheduled_departure,scheduled_arrival,departure_airport,arrival_airport,status,aircraft_code,actual_departure,actual_arrival,seat_no,fare_conditions
2 1185,PG0134,2017-09-10 09:50:00+03,2017-09-10 14:55:00+03,DME,BTK,Scheduled,319,,
3 3979,PG0052,2017-08-25 14:50:00+03,2017-08-25 17:35:00+03,VKO,HMA,Scheduled,CR2,,
4 4739,PG0561,2017-09-05 12:30:00+03,2017-09-05 14:15:00+03,VKO,AER,Scheduled,763,,
5 5502,PG0529,2017-09-12 09:50:00+03,2017-09-12 11:20:00+03,SVO,UFA,Scheduled,763,,
6 6938,PG0461,2017-09-04 12:25:00+03,2017-09-04 13:20:00+03,SVO,ULV,Scheduled,SU9,,
7 7784,PG0667,2017-09-10 15:00:00+03,2017-09-10 17:30:00+03,SVO,KRO,Scheduled,CR2,,
8 9478,PG0360,2017-08-28 09:00:00+03,2017-08-28 11:35:00+03,LED,REN,Scheduled,CR2,,
9 11085,PG0569,2017-08-24 15:05:00+03,2017-08-24 16:10:00+03,SVX,SCW,Scheduled,733,,
10 11847,PG0498,2017-09-12 10:15:00+03,2017-09-12 14:55:00+03,KZN,IKT,Scheduled,319,,
11 12012,PG0621,2017-08-26 16:05:00+03,2017-08-26 17:00:00+03,KZN,MQF,Scheduled,CR2,,
12 13113,PG0612,2017-08-18 16:25:00+03,2017-08-18 20:05:00+03,ROV,KZN,Scheduled,CN1,,
13 14806,PG0676,2017-09-06 07:05:00+03,2017-09-06 07:45:00+03,PEE,CEK,Scheduled,CR2,,
14 16837,PG0010,2017-09-05 12:25:00+03,2017-09-05 14:35:00+03,JOK,VKO,Scheduled,CN1,,
15 17173,PG0059,2017-09-14 12:25:00+03,2017-09-14 14:45:00+03,SCW,NBC,Cancelled,CN1,,
16 19807,PG0035,2017-09-11 06:35:00+03,2017-09-11 09:25:00+03,MJZ,CNN,Scheduled,CN1,,
17 23609,PG0648,2017-08-31 11:35:00+03,2017-08-31 13:00:00+03,UUA,SVO,Scheduled,CR2,,
18 23695,PG0388,2017-08-26 10:55:00+03,2017-08-26 11:25:00+03,UUA,REN,Scheduled,CR2,,
19 23780,PG0098,2017-09-02 06:50:00+03,2017-09-02 10:30:00+03,SWT,CEK,Scheduled,CN1,,
20 23945,PG0076,2017-09-05 09:15:00+03,2017-09-05 11:50:00+03,EYK,DME,Scheduled,CR2,,
21 24705,PG0632,2017-08-26 15:00:00+03,2017-08-26 17:35:00+03,TJM,PES,Scheduled,CR2,,
22 25382,PG0556,2017-08-31 09:05:00+03,2017-08-31 10:55:00+03,IKT,MJZ,Scheduled,CR2,,
23 26057,PG0467,2017-08-26 15:45:00+03,2017-08-26 17:00:00+03,ULY,RTW,Scheduled,CN1,,
24 27580,PG0483,2017-09-12 07:20:00+03,2017-09-12 11:20:00+03,KEJ,DME,Scheduled,SU9,,
25 29272,PG0334,2017-09-05 16:20:00+03,2017-09-05 19:05:00+03,KGD,ESL,Scheduled,CR2,,
26 29440,PG0065,2017-09-02 12:15:00+03,2017-09-02 18:05:00+03,UUD,VKO,Scheduled,319,,
27 32658,PG0674,2017-08-19 09:35:00+03,2017-08-19 12:00:00+03,KRO,KJA,Scheduled,CR2,,
28 1,PG0405,2017-07-16 09:35:00+03,2017-07-16 10:30:00+03,DME,LED,Arrived,321,2017-07-16 09:44:00+03,2017-07-16 10:39:00+03,,
29 2,PG0404,2017-08-05 19:05:00+03,2017-08-05 20:00:00+03,DME,LED,Arrived,321,2017-08-05 19:06:00+03,2017-08-05 20:01:00+03,,
30 3,PG0405,2017-08-05 09:35:00+03,2017-08-05 10:30:00+03,DME,LED,Arrived,321,2017-08-05 09:39:00+03,2017-08-05 10:34:00+03,,
```

AI Model 2 – Real-World Data : Train/test splitting and Balancing Classes

```
38 # --- Train/test split ---
39
40 X_train, X_test, y_train, y_test = train_test_split(
41     X, y, test_size=0.3, random_state=42, stratify=y)
42
43 # --- Handle class imbalance by undersampling majority class ---
44
45 # Combine X_train and y_train for easier resampling
46 train_data = pd.concat([X_train, y_train], axis=1)
47
48 # Separate majority and minority classes
49 majority = train_data[train_data.status_Delayed == 0]
50 minority = train_data[train_data.status_Delayed == 1]
51
52 print("Before undersampling:", majority.status_Delayed.value_counts(), minority.status_Delayed.value_counts())
53
54 # Downsample majority class
55 majority_downsampled = resample(
56     majority,
57     replace=False,
58     n_samples=len(minority),
59     random_state=42
60 )
61
62 # Combine minority class with downsampled majority class
63 undersampled = pd.concat([majority_downsampled, minority])
64
65 print("After undersampling:", undersampled.status_Delayed.value_counts())
66
67 # Split back into X and y
68 X_train_bal = undersampled.drop('status_Delayed', axis=1)
```

Applied resampling to address class imbalance

AI Model 2 – Real-World Data : Training Model and Accuracy

```
71 # --- Train RandomForestClassifier model ---
72
73 clf = RandomForestClassifier(random_state=42)
74 clf.fit(X_train_bal, y_train_bal)
75
76 # --- Evaluate model ---
77
78 y_pred = clf.predict(X_test)
79
80 print("Accuracy:", accuracy_score(y_test, y_pred))
81 print(classification_report(y_test, y_pred, digits=4))
82
```

RandomForestClassifier achieved 92% accuracy but F1 score for delays remained low (0.03), highlighting real-world challenges

```
Before undersampling: status_Delayed
0    24093
Name: count, dtype: int64 status_Delayed
1      29
Name: count, dtype: int64
After undersampling: status_Delayed
0      29
1      29
Name: count, dtype: int64
Accuracy: 0.9235828980460438
```

	precision	recall	f1-score	support
0	1.0000	0.9235	0.9602	10326
1	0.0150	1.0000	0.0295	12
accuracy			0.9236	10338
macro avg	0.5075	0.9617	0.4949	10338
weighted avg	0.9989	0.9236	0.9591	10338

Output Results

F1 score for "Not Delayed" class: 0.9602

F1 score for "Delayed" class: 0.0295


AI Model 2 – Real-World Data : Top 10 Important Features

```
[5]: 1 # --- Feature Importance from trained model ---
2 importances = clf.feature_importances_
3 feature_names = X_train.columns
4
5 # Create DataFrame of features and their importances
6 feature_importance_df = pd.DataFrame({
7     'Feature': feature_names,
8     'Importance': importances
9 }).sort_values(by='Importance', ascending=False)
10
11 # Print top 10 most important features
12 print("Top 10 Important Features:")
13 print(feature_importance_df.head(10))
14
15 # --- Plotting ---
16 plt.figure(figsize=(10, 6))
17 sns.barplot(
18     data=feature_importance_df.head(10),
19     x='Importance',
20     y='Feature',
21     hue='Feature',
22     dodge=False,
23     palette='viridis',
24     legend=False
25 )
26 plt.title('Top 10 Feature Importances')
27 plt.xlabel('Importance Score')
28 plt.ylabel('Feature')
29 plt.tight_layout()
30 plt.show()
```

Top 10 Important Features:

	Feature	Importance
3	scheduled_arrival	0.366114
2	scheduled_departure	0.346360
7	actual_departure	0.077389
5	arrival_airport	0.043414
0	flight_id	0.040492
8	actual_arrival	0.034601
1	flight_no	0.028396
4	departure_airport	0.023876
6	aircraft_code	0.018329
9	seat_no	0.017972

70% delays from
scheduling



Reflection

- The limitations of synthetic vs. real-world datasets.
- The difficulty of generalizing models to imbalanced, messy real data.
- Building, debugging, and refining the engineering pipeline mattered as much as accuracy itself.

Demo and Project links

- **GitHub Repo:** <https://github.com/Gordonandric/Flight-Delay-Prediction-Modeling>
- **Kaggle:**
 - <https://www.kaggle.com/writeups/gordonandric/flight-delay-prediction-modeling>
 - <https://www.kaggle.com/code/gordonandric/airline-synthetic>
 - <https://www.kaggle.com/datasets/gordonandric/synthetic-airline-data>
 - <https://www.kaggle.com/code/gordonandric/airline-real>
 - <https://www.kaggle.com/datasets/gordonandric/real-airline-data/>
- **Demo:**
 - https://github.com/Gordonandric/Flight-Delay-Prediction-Modeling/blob/main/Airline_Synthetic.webm
 - https://github.com/Gordonandric/Flight-Delay-Prediction-Modeling/blob/main/Airline_Real.webm

Thank You